

MASHWARA

Enterprise Product Requirements Document

Executive Intelligence Twin for Pakistan

State memory, live meeting intelligence, user-created agents, and execution accountability for PM and ministerial workflows.

Version v1.0 Enterprise Draft | May 3, 2026

Classification: Confidential Draft

Document Control

Field	Value
Product	Mashwara
Document type	Enterprise Product Requirements Document (PRD)
Version	v1.0 Enterprise Draft
Date	May 3, 2026
Primary audience	Product, Engineering, Design, Data/AI, Security, Program Leadership, PMO/Ministerial stakeholders
Working classification	Confidential Draft
Source baseline	Based on the uploaded Wakeel mobile PRD and subsequent Mashwara requirements discussed in this conversation. The enterprise version replaces the original persona, data model, and scope with Pakistan-specific RAG, live meeting, agent, execution, and enterprise-governance requirements.

Table of Contents

- 1. Executive Summary
- 2. Requirement Review and Pakistan Shift
- 3. Product Vision and Enterprise Positioning
- 4. Problem Statement and Outcomes
- 5. Personas, Roles, and Access Model
- 6. Product Principles
- 7. Product Scope and Core Modules
- 8. Data Source and RAG Layers
- 9. User Stories
- 10. Functional Requirements
- 11. Non-Functional Requirements
- 12. Conceptual Data Model
- 13. Key User Journeys
- 14. Enterprise Rollout Plan
- 15. Success Metrics
- 16. Risks, Assumptions, and Open Questions
- 17. Out of Scope
- Appendix A. Requirement Index

1. Executive Summary

Mashwara is an enterprise executive intelligence system for Pakistan federal leadership. It transforms fragmented **official communication**, **calendars**, **minutes**, **documents**, **live meeting recordings**, and **execution logs** into a source-backed operating layer for decisions, meetings, agents, and accountability.

The product is designed for Prime Minister, Federal Minister, Secretary, PM Office/private office, and ministry execution workflows. It is not just a briefing application. It is a real-time and historical intelligence twin that can discover decisions from existing sources, record live meetings on mobile, run agents over live or historical context, create user-defined agents, and track execution until closure.

Enterprise v1 expands the baseline mobile executive experience into a governed, multi-role, permission-aware platform. The core enterprise requirement is that every generated output must be explainable, source-linked, permission-scoped, auditable, and actionable.

2. Requirement Review and Pakistan Shift

The original product pattern provides useful executive surfaces: brief, decisions, meetings, Ask, agents/signals, and decision history. For Pakistan, the requirement must shift from a single-executive mobile assistant to an enterprise system with federal governance, live meeting capture, user-created agents, RAG-first memory, and execution accountability.

Area	Enterprise PRD decision
Product name	Use Mashwara as the Pakistan product name. Position it as trusted counsel, institutional memory, and real-time decision intelligence.
Persona model	Move from a single apex executive persona to a Pakistan federal executive model: Prime Minister, Federal Ministers, Secretaries, PM Office / private office staff, ministry focal points, and enterprise administrators.
Data posture	Move from curated app content to RAG-first discovery across Outlook, calendar, minutes of meeting, pre-reads, documents, live meeting transcripts, and execution logs.
Meeting posture	Add mobile Start Meeting capture for scheduled and ad-hoc meetings, with live transcription, reviewable minutes, decision/action extraction, and source-linked transcript evidence.
Agent posture	Upgrade agents from browse/manage signals to user-created agents that can run on historical RAG, upcoming meetings, live meetings, decisions, and execution logs.
Enterprise posture	Add tenant administration, RBAC/ABAC, permission-aware retrieval, audit trails, retention policy, consent policy, connector operations, and compliance controls.
Decision posture	Shift from decision workflow only to decision + execution lifecycle: detect, recommend, approve, issue directive, assign owner, monitor delivery, flag delays, and close with evidence.

3. Product Vision and Enterprise Positioning

Product name: Mashwara.

Positioning statement: Mashwara is the institutional memory and real-time decision intelligence layer for Pakistan federal executive workflows.

Tagline: Trusted counsel for state decisions, live meetings, agents, and execution.

3.1 Product Thesis

- Pakistan executive workflows generate decisions in many places: emails, calendar meetings, minutes, informal reviews, live discussions, documents, and follow-up messages.
- Mashwara should not wait for a user to manually create work. It should discover decision candidates, risks, commitments, and delays from reality.
- The product must capture the meeting while it happens, then convert that meeting into reviewable minutes, decisions, actions, and execution items.
- Users must be able to create their own agents to monitor issues, run over meetings, track execution, and surface signals without engineering intervention.

4. Problem Statement and Outcomes

4.1 Problems to Solve

Problem	Description
Fragmented inputs	Emails, meetings, documents, and minutes are spread across systems, folders, staff, and ministries.
Institutional amnesia	Important commitments and earlier decisions are difficult to retrieve during new meetings and policy reviews.
Weak execution trail	Decisions may be made, but ownership, due dates, evidence, and closure are often tracked separately or manually.
Inter-ministry misalignment	Different ministries may hold conflicting positions without a single view of disagreement and impact.
Uncaptured live decisions	Important decisions are made in meetings before minutes are drafted, delayed, or corrected.
Limited personalized monitoring	Leaders need issue-specific monitors, but cannot wait for engineering teams to configure every new tracker.

4.2 Enterprise Outcomes

- Officials start the day with a complete, sourced, and role-appropriate brief.
- Decisions are discovered automatically, explained with evidence, approved by humans, and tracked through execution.
- Meetings generate prep packs before and reviewable MoMs after, while live meetings can be captured from mobile.
- Agents can be created by authorized users and run on emails, meetings, decisions, transcripts, and execution logs.
- Every insight, recommendation, directive, and action item has source provenance and auditability.
- Enterprise administrators can govern access, connectors, retention, recording policies, agent behavior, and audits.

5. Personas, Roles, and Access Model

Persona	Role	Primary needs
Prime Minister	Primary executive decision maker across ministries and national priorities.	Cross-ministry brief, high-stakes decisions, national execution status, crisis updates, and issue lineage.
Federal Minister	Domain executive for one or more portfolios.	Ministry brief, policy decisions, meeting prep, execution follow-up, and issue-specific agents.
Secretary	Administrative lead responsible for execution and formal records.	Action ownership, minutes review, directive acknowledgement, execution evidence, and issue escalation.
Principal Secretary / Chief of Staff	Executive office operator and gatekeeper.	Prepare PM/Minister day, review brief items, start or supervise meetings, triage decision queue, manage sensitive agents.
Secretariat / Private Office Staff	Operational support for calendar, records, correspondence, and follow-ups.	Correct extracted actions, validate MoMs, manage notifications, route directives, and maintain source hygiene.
Ministry Focal Point	Designated official for a policy area or execution item.	Receive assigned actions, provide updates, attach evidence, and respond to agent-generated questions.
Enterprise Admin	IT/security/operator role without default content access.	Manage tenants, connectors, identity, policies, observability, retention, and support workflows.
Compliance / Audit Viewer	Read-only oversight role.	Review audit logs, evidence chains, retention compliance, and access history under authorized scope.

5.1 Access Model Principles

- Prime Minister and PM Office roles may require cross-ministry visibility, but only under explicit policy and approved source scopes.
- Federal Ministers should see their portfolio/ministry scope and specifically delegated cross-ministry items.
- Secretaries and focal points should see execution and source material relevant to assigned items, not unrelated executive content.
- Enterprise admins should operate configuration and health without default access to confidential content.
- Agents inherit creator permissions and configured scope. Agents cannot expand access beyond authorized sources.
- Every access, generation, export, edit, recording, and agent run must produce audit events.

6. Product Principles

- Reality first: Mashwara discovers work from actual communication, meetings, documents, and live conversations rather than relying on manual entry.
- Human authority: Mashwara recommends, drafts, and tracks; the authorized official remains the final decision maker.
- Source-visible by default: every answer, recommendation, decision, and action item must show the evidence that produced it.
- Permission-aware intelligence: retrieval, summaries, agents, and search must honor source-system permissions and Mashwara roles.
- Live plus historical: the system must understand both the real-time meeting happening now and the institutional memory that came before it.
- Execution over activity: the product is successful only when decisions move into accountable execution and close with evidence.
- Configurable sensitivity: recording, agent behavior, retention, redaction, and sharing must be policy-controlled for sensitive government contexts.

7. Product Scope and Core Modules

Enterprise v1 should include both executive-facing mobile experiences and enterprise administration controls. Mobile is the primary executive surface; an admin web console is required for enterprise setup and governance.

Module	Purpose	Primary surfaces
National Brief	Daily executive brief generated from RAG, live meeting carryovers, decisions, execution status, and upcoming calendar.	Brief tab, audio brief, item drilldown, high-stakes notification.
Decision Engine	Extracts, prioritizes, explains, and routes decisions from emails, meetings, documents, agents, and live sessions.	Decision queue, decision detail, reasoning, directive draft, receipt, history.
Meetings Intelligence	Prepares officials before meetings and converts meetings into summaries, decisions, actions, and continuity records.	Calendar timeline, prep pack, post-summary, continuity layer, document preview.
Live Meeting Capture	Lets authorized users start a scheduled or ad-hoc meeting recording on mobile and run agents during the meeting.	Start Meeting, live transcript, live agents, detected decisions/actions, review and publish.
Agent Studio	Allows users to create, run, manage, and review agents using natural language and approved templates.	Create Agent, Your Agents, Running Now, Agent Detail, Agent Outputs.
Execution Monitor	Tracks decision execution from directive to owner assignment, status update, delay detection, escalation, and closure.	Execution dashboard, action detail, ministry view, overdue flags, evidence closure.
Ask Mashwara	A permission-aware intelligence interface across the full memory graph with citations and confidence.	Global Ask, voice input, file evaluation, context-aware suggestions, conversation history.
Enterprise Admin	Controls identity, connectors, policies, retention, audit, data quality, tenant configuration, and operational observability.	Admin console, connector dashboard, role manager, audit viewer, policy manager.

7.1 Priority Definitions

Priority	Definition
P0	Launch blocker for enterprise pilot; must be implemented or explicitly waived before production use.
P1	Required for a credible enterprise MVP; can be phased during pilot if a manual workaround is accepted.
P2	Enhancement or scale feature; not required for initial pilot but needed for full national rollout.

8. Data Source and RAG Layers

Mashwara is RAG-first and live-first. The system must understand historical evidence, future meeting context, and live meeting streams under a single permission-aware memory model.

Layer	Inputs	Primary outputs
Outlook RAG	Official emails, threads, attachments, sender/recipient metadata, timestamps, forwarded context, and follow-up patterns.	Decision candidates, escalations, stakeholder positions, missing replies, external completion signals.
Calendar RAG	Scheduled meetings, attendees, agenda titles, recurrence, location, pre-read links, and organizer information.	Meeting timeline, prep packs, participant context, expected decisions, and upcoming risk flags.

Layer	Inputs	Primary outputs
MoM RAG	Historical minutes, cabinet/committee records, secretary-level minutes, decisions, commitments, and unresolved actions.	Institutional memory, decision lineage, carryovers, unresolved items, and accountability trail.
Upcoming Meetings RAG	Future agendas, pre-reads, linked documents, stakeholder lists, and expected decision topics.	Pre-meeting briefing, talking points, issue history, conflict signals, and agent pre-runs.
Live Meeting Capture	Mobile audio recording, transcript, speaker segments, tags, and in-meeting agent observations.	Draft MoM, decisions, action items, commitments, conflicts, and transcript-backed evidence.
Execution Logs	Directive receipts, assignments, status updates, comments, evidence attachments, and completion records.	Delivery monitoring, delay alerts, silent failure detection, escalation, and closure proof.
Document Repositories	Policy papers, memos, presentations, reports, spreadsheets, and scanned records where connectors are approved.	Source packs, document summaries, proposal evaluation, and reference evidence.

8.1 Conceptual Data Flow

Step	Description
1. Connect	Approved connectors ingest Outlook, Calendar, MoM repositories, document repositories, live audio/transcripts, and execution systems.
2. Normalize	Content is parsed, deduplicated, classified, permission-tagged, and linked to source provenance.
3. Retrieve	RAG retrieves only authorized context for each user, surface, query, decision, or agent.
4. Generate	Mashwara generates briefs, recommendations, prep packs, summaries, agent outputs, and answers with source references.
5. Review	Human reviewers correct/approve sensitive or official artifacts before dispatch/publication.
6. Execute	Approved decisions create directives and execution items; progress is monitored until closure.
7. Audit	Every important event is logged for authorized review and compliance.

9. User Stories

ID	As a	I want to	So that
US-BRF-001	Prime Minister	receive a daily national brief with critical decisions, execution delays, conflicts, and high-stakes meetings	start the day with an actionable state view
US-BRF-002	Federal Minister	listen to an audio version of the brief	prepare while in transit
US-DEC-001	Prime Minister or Minister	open a decision and see recommendation, confidence, source evidence, ministries involved, and lineage	decide without chasing background
US-DEC-002	authorized official	edit and send a directive drafted from a decision	retain authorship while reducing manual drafting time
US-MTG-001	Minister	open a prep pack before a scheduled meeting	walk in with agenda, history, stakeholder stance, and unresolved actions
US-LIVE-001	authorized mobile user	tap Start Meeting and record an ad-hoc or scheduled meeting	capture decisions even when no formal MoM process exists yet
US-LIVE-002	meeting lead	run selected agents during the meeting	detect commitments, disagreements, risks, and follow-ups in real time
US-LIVE-003	secretariat reviewer	review, correct, and publish AI-generated minutes after the meeting	formalize a reliable record before it enters execution
US-AGT-001	Prime Minister or Minister	create an agent in natural language	monitor issues without writing configuration
US-AGT-002	Secretary	run an agent on a specific meeting or policy issue	surface risks and missing commitments before execution starts
US-EXE-001	Secretary	see all assigned actions with due dates and evidence requirements	coordinate implementation and report progress
US-ASK-001	authorized user	ask questions across emails, meetings, decisions, transcripts, and execution logs	retrieve institutional memory with source-backed answers
US-ADM-001	enterprise admin	manage connectors, roles, retention, and audit policies	operate Mashwara safely at enterprise scale

10. Functional Requirements

Requirement IDs are grouped by product capability. Acceptance criteria are written at product level; engineering should later decompose them into technical tasks and test cases.

App Shell, Identity, and Navigation

ID	Priority	Requirement	Acceptance Criteria
FR-APP-001	P0	Authenticate users through enterprise SSO with MFA and optional biometric session unlock on mobile.	A user cannot access Mashwara without a valid enterprise identity; mobile re-entry can be gated by biometric/PIN based on policy.
FR-APP-002	P0	Render navigation based on role and authorized modules: Brief, Decisions, Meetings, Live, Agents, Execution, Ask, and Settings/Admin as applicable.	A Minister sees only authorized portfolio data; an admin can manage system settings without default access to confidential content.
FR-APP-003	P0	Expose Ask Mashwara globally as a persistent, context-aware control across core user screens.	Opening Ask from a decision, meeting, transcript, or execution item carries the current context into the question panel.
FR-APP-004	P1	Support secure push notifications for prep packs, meeting summaries, high-priority decisions, agent alerts, and overdue execution items.	Notification text must avoid sensitive details when the device is locked; tapping the notification deep-links after authentication.
FR-APP-005	P1	Allow user preferences for brief time, notification intensity, default language, audio voice, and sensitive-mode defaults.	Preferences persist per user and apply to generated brief delivery and mobile capture behavior.
FR-APP-006	P1	Provide offline read access to the latest downloaded brief, meeting prep pack, and cached decision detail, excluding send/execute actions.	When offline, read-only views are clearly marked; sending directives and syncing meeting recordings are disabled until connected.

Data Ingestion and RAG Foundation

ID	Priority	Requirement	Acceptance Criteria
FR-RAG-001	P0	Ingest Outlook email threads and attachments from approved mailboxes and distribution groups.	Email threads, sender/recipient metadata, timestamps, attachments, and replies are indexed with source links and permission metadata.
FR-RAG-002	P0	Ingest official calendars for authorized PM/Minister/Secretary accounts and selected shared calendars.	Scheduled meetings appear in Mashwara with participants, agenda metadata, linked documents, and pre/post meeting status.
FR-RAG-003	P0	Ingest historical minutes of meeting repositories and formal decision records.	Past decisions, action items, owners, dates, and meeting references can be retrieved as evidence in brief, decision, and Ask outputs.
FR-RAG-004	P0	Ingest upcoming meeting agendas and pre-read material before meeting time.	Prep packs include agenda, documents, stakeholders, issue history, prior decisions, and unresolved actions.
FR-RAG-005	P0	Normalize and classify unstructured content into source document, email thread, meeting, decision candidate, action item, stakeholder position, and evidence objects.	The system can show how a decision candidate was created from one or more source objects.
FR-RAG-006	P0	Apply permission-aware retrieval for every RAG response, brief item, decision, agent output, and search result.	A user cannot retrieve or infer content from a source they are not authorized to access.
FR-RAG-007	P0	Maintain durable provenance for every generated insight.	Every generated brief item, recommendation, summary, answer, and agent output links back to one or more source records.
FR-RAG-008	P1	Detect duplicates, near-duplicates, thread continuations, and renamed documents across sources.	The same issue is not shown as multiple unrelated decisions unless the system explains the distinction.
FR-RAG-009	P1	Capture human corrections to improve future extraction and classification.	When a reviewer changes an owner, deadline, decision type, or summary, the correction is stored as feedback with audit metadata.

National Brief

ID	Priority	Requirement	Acceptance Criteria
FR-BRF-001	P0	Generate a daily National Brief for each executive based on authorized sources and role scope.	The brief is generated by configured time and contains only authorized content.
FR-BRF-002	P0	Include required sections: Critical Decisions Today, Pending Since Last Cabinet/Committee, Execution Delays, Cross-Ministry Conflicts, Upcoming High-Stakes Meetings, and Agent Alerts.	Each section can be empty, but an empty state explains why no items are shown.
FR-BRF-003	P0	Support a 3-5 minute audio brief with chapter navigation.	The user can play, pause, scrub, and jump to chapters; each audio chapter maps to visible brief sections.
FR-BRF-004	P0	Make every brief item drillable into the underlying decision, meeting, source pack, execution item, or agent output.	Tapping a brief item opens the most relevant detail view without manual search.
FR-BRF-005	P1	Highlight carryovers and stale items based on prior meetings, unresolved decisions, and execution status.	Items older than configured thresholds show age, source, owner, and why they are still relevant.
FR-BRF-006	P1	Allow the private office or authorized reviewer to mark brief items as sensitive, suppress from audio, or require review before publication.	Sensitive items obey policy and display review status.

Decision Engine

ID	Priority	Requirement	Acceptance Criteria
FR-DEC-001	P0	Automatically extract decision candidates from emails, meetings, documents, live transcripts, and agent outputs.	A decision candidate contains title, source, reason for extraction, urgency, involved ministries, and recommended next step.
FR-DEC-002	P0	Support decision lanes: Needs Review, Ready, Waiting on Entity, Executing, Deferred, Completed, and Closed Without Action.	Lane changes are server-authoritative, audit logged, and visible in decision history.
FR-DEC-003	P0	Display decision detail with at-a-glance summary, recommendation, confidence, ministries involved, stakeholder positions, source pack, related meetings, related decisions, and execution status.	A user can understand context without leaving the decision detail screen.
FR-DEC-004	P0	Generate one of four primary recommendations: Approve, Approve with Conditions, Return for Clarification, or Reject.	The recommendation includes reasoning, confidence, criteria, and evidence links.
FR-DEC-005	P0	Provide a one-tap action path into an editable directive draft.	The user can open the draft from a decision action, edit To/CC/Subject/Body, and cancel without losing decision context.
FR-DEC-006	P0	Require explicit human confirmation before sending or registering any directive.	No directive is dispatched solely by AI or agent action.
FR-DEC-007	P0	Generate a registered decision receipt after send/approval with reference, timestamp, action, recipients, and status row.	The receipt is stored and searchable in decision history.
FR-DEC-008	P0	Maintain decision lineage from first mention to current state.	A decision shows when it first appeared, meetings where it was discussed, emails/documents linked, prior recommendations, and what changed.
FR-DEC-009	P1	Detect potential external completion from email replies, signed documents, or execution evidence.	The system can mark a decision as possibly completed externally and require reviewer confirmation.
FR-DEC-010	P0	Maintain a full audit trail of decision state, source changes, reviewer edits, approvals, deferrals, and dispatch events.	Audit events are immutable to normal users and available to authorized audit roles.
FR-DEC-011	P1	Detect conflicting ministry positions and attach a conflict signal to the decision.	Decision detail shows ministries, disagreement summary, source evidence, and suggested reconciliation questions.
FR-DEC-012	P2	Suggest precedent decisions and similar historical cases.	Related decisions are ranked by issue, ministries, policy area, and outcome similarity.

Meetings Intelligence

ID	Priority	Requirement	Acceptance Criteria
FR-MTG-001	P0	Show a role-aware meeting timeline with scheduled, in-progress, completed, and ad-hoc live sessions.	A user can navigate today, past days, and upcoming days based on calendar permission.
FR-MTG-002	P0	Generate a prep pack for each upcoming meeting with context, agenda, issue history, stakeholders, prior decisions, unresolved actions, documents, and suggested talking points.	Prep packs are source-backed and available from notifications, calendar rows, and brief links.
FR-MTG-003	P0	Generate post-meeting summaries containing summary, decisions, action items, commitments, follow-ups, owners, due dates, and source evidence.	Post-summary is reviewable, editable by authorized reviewers, and can create decisions/actions.
FR-MTG-004	P0	Provide a continuity layer for recurring or related meetings.	The meeting view shows unresolved actions from prior meetings and what changed since the previous session.
FR-MTG-005	P1	Support in-app document preview for approved formats and source packs.	Users can view documents without switching tools; restricted documents remain inaccessible.
FR-MTG-006	P1	Send configurable prep and summary notifications.	Prep notifications default to a configurable pre-meeting window; summary notifications appear after summary generation/review policy.
FR-MTG-007	P0	Link meeting outcomes to Decision Engine and Execution Monitor.	A decision or action extracted from a meeting appears in the appropriate queue with meeting evidence attached.

Live Meeting Capture

ID	Priority	Requirement	Acceptance Criteria
FR-LIVE-001	P0	Allow authorized users to tap Start Meeting on mobile for scheduled or ad-hoc meetings.	The user can start capture from the Meetings tab, Brief, notification, or floating action button.
FR-LIVE-002	P0	Allow the user to attach the live session to a calendar meeting or create an ad-hoc meeting record.	The live session has title, time, participants, organizer, sensitivity level, and optional agenda.
FR-LIVE-003	P0	Enforce configurable recording consent and sensitivity policy before capture begins.	The start flow displays required consent/sensitivity prompts and blocks recording if policy conditions are not met.
FR-LIVE-004	P0	Capture mobile audio reliably with clear recording status and network/offline state.	The user sees elapsed time, recording state, sync state, and warnings for low battery/network issues.
FR-LIVE-005	P0	Generate live or near-live transcription with speaker segmentation where supported.	Transcript segments include timestamp, speaker label or unknown speaker, confidence, and source link to audio segment.
FR-LIVE-006	P0	Detect decisions, commitments, action items, risks, disagreements, deadlines, and named owners during the meeting.	Detected items appear in a reviewable live panel and are not finalized until accepted by an authorized reviewer.
FR-LIVE-007	P0	Allow users to run selected agents on top of the live	At least one agent can run during capture and generate live

ID	Priority	Requirement	Acceptance Criteria
		meeting.	observations without interrupting recording.
FR-LIVE-008	P1	Support in-meeting controls: pause recording, mark off-record segment, bookmark moment, tag topic, and flag sensitive segment.	Controls create timestamped audit events and affect transcript visibility according to policy.
FR-LIVE-009	P0	Generate a review package when the meeting ends.	The package includes transcript, draft MoM, detected decisions/actions, agent outputs, unresolved questions, and source evidence.
FR-LIVE-010	P0	Route accepted decisions and actions into Decision Engine and Execution Monitor.	Accepted items appear in the right queues with meeting transcript evidence and reviewer attribution.
FR-LIVE-011	P1	Allow post-meeting correction of transcript, speaker labels, summaries, owners, and deadlines.	Corrections are audit logged and update downstream summaries/actions where appropriate.
FR-LIVE-012	P2	Support multi-device or room-system capture when enterprise hardware is available.	A meeting can merge mobile capture with approved room audio while maintaining a single transcript and audit record.

Agent Studio and Agent Execution

ID	Priority	Requirement	Acceptance Criteria
FR-AGT-001	P0	Allow authorized users to create agents using natural language.	The user can describe what to monitor or analyze; Mashwara converts it into a draft agent configuration for review.
FR-AGT-002	P0	Provide approved agent templates for monitoring, meeting capture, decision review, conflict detection, execution tracking, and weekly reporting.	Users can start from a template and edit objective, sources, triggers, outputs, and recipients.
FR-AGT-003	P0	Allow agent configuration of data scope, triggers, schedule, output type, recipients, and sensitivity.	The system validates configuration against permissions and policy before activation.
FR-AGT-004	P0	Allow agents to run on historical RAG sources, upcoming meetings, a selected decision, a live meeting, or execution logs.	Agent run context is visible and stored with each output.
FR-AGT-005	P0	Support real-time meeting agents that detect commitments, contradictions, missing owners, policy references, and follow-up questions.	During a live meeting, agent observations show timestamp, confidence, and transcript/source evidence.
FR-AGT-006	P0	Support agent output types: signal, alert, report, suggested question, draft action, and decision candidate.	Output type controls routing to Brief, Decisions, Execution, Notifications, or Agent Inbox.
FR-AGT-007	P0	Allow users to pause, resume, run now, edit, duplicate, deactivate, and view run history for their agents.	Every state change is audit logged; edits create a new version.
FR-AGT-008	P0	Apply guardrails: agents cannot access unauthorized sources, cannot dispatch directives autonomously, and cannot override sensitivity policy.	Policy violations block activation or run with a clear explanation.
FR-AGT-009	P1	Provide test-run preview before activating an agent.	A test run shows sample outputs and sources without sending notifications or creating decisions.
FR-AGT-010	P1	Support shared agents managed by a private office or ministry team.	Shared agents show owner, scope, subscribers, and approval status.
FR-AGT-011	P2	Provide an enterprise template library governed by admins.	Admins can publish, retire, and version approved templates.

Execution Monitor

ID	Priority	Requirement	Acceptance Criteria
FR-EXE-001	P0	Create execution items from directives, MoMs, live meeting actions, and accepted agent outputs.	Each item contains owner, ministry/entity, due date, source, status, priority, and evidence requirement.
FR-EXE-002	P0	Track status across Not Started, In Progress, At Risk, Blocked, Completed, Deferred, and Cancelled.	Status changes show timestamp, updater, source, comment, and optional evidence.
FR-EXE-003	P0	Detect overdue actions and silent failures.	Items with no update after a configured period are flagged and surfaced in Brief/Execution/Notifications.
FR-EXE-004	P1	Support escalation from ministry focal point to Secretary, Minister, or PM Office based on policy.	Escalations are logged and visible in the execution timeline.
FR-EXE-005	P0	Allow closure only with completion evidence or authorized override.	Completion requires evidence attachment, source link, or explicit override reason.
FR-EXE-006	P1	Provide ministry and cross-ministry execution dashboards.	Dashboards summarize overdue, at-risk, completed, and blocked items by ministry, policy area, and meeting/decision source.
FR-EXE-007	P1	Generate weekly execution summaries for authorized leaders.	Summaries include progress, delays, blockers, and recommended escalations with evidence links.
FR-EXE-008	P2	Support integration with external workflow or document-signature systems where approved.	External status updates can update execution items with source provenance.

Ask Mashwara

ID	Priority	Requirement	Acceptance Criteria
FR-ASK-001	P0	Allow users to query across authorized emails, meetings, MoMs, transcripts, decisions, agents, documents, and execution logs.	Answers are scoped to user permissions and current context.
FR-ASK-002	P0	Show source citations and confidence for every factual	The user can open source evidence from the answer.

ID	Priority	Requirement	Acceptance Criteria
		answer.	
FR-ASK-003	P0	Clearly state uncertainty, missing data, or insufficient authorization.	The system does not fabricate answers when evidence is unavailable.
FR-ASK-004	P0	Provide context-aware suggestion chips on key screens.	On a meeting, suggested questions relate to prep/summary; on a decision, suggestions relate to reasoning/lineage/execution.
FR-ASK-005	P1	Support voice input and voice playback where policy allows.	Voice interactions are transcribed, permission-scoped, and stored according to retention policy.
FR-ASK-006	P1	Allow users to upload or reference a file for evaluation.	Mashwara summarizes, evaluates, and links outputs to the file with policy-compliant retention.
FR-ASK-007	P1	Support conversation history with topic, source scope, and continuation.	Users can resume prior conversations without re-uploading context, subject to retention settings.
FR-ASK-008	P0	Block or safely decline requests that would violate role permissions, data policy, or autonomous decision boundaries.	The refusal explains the policy reason without exposing sensitive hidden content.

Enterprise Administration and Governance

ID	Priority	Requirement	Acceptance Criteria
FR-ADM-001	P0	Provide tenant and workspace configuration for PM Office and ministries.	Admins can configure organization units, ministry scopes, roles, and feature flags.
FR-ADM-002	P0	Provide RBAC/ABAC policy management.	Admins can define access by role, ministry, clearance, meeting sensitivity, source system, and explicit delegation.
FR-ADM-003	P0	Provide connector management for Outlook, calendars, MoM repositories, document repositories, and execution systems.	Admins can enable, disable, test, monitor, and rotate credentials for approved connectors.
FR-ADM-004	P0	Provide retention and deletion policy controls for transcripts, audio, summaries, agent outputs, decisions, and audit logs.	Retention policy is enforceable per source type, sensitivity, and legal/records category.
FR-ADM-005	P0	Provide audit logs for access, generation, exports, edits, approvals, agent runs, recording events, and policy changes.	Authorized auditors can filter and export logs based on permitted scope.
FR-ADM-006	P0	Provide recording consent and sensitivity policy configuration.	Admins can define when recording is allowed, what prompts are required, and who can access audio/transcripts.
FR-ADM-007	P1	Provide data quality and ingestion health dashboards.	Admins can see connector status, failed ingestions, stale indexes, OCR/parse failures, and source coverage.
FR-ADM-008	P1	Provide model, prompt, and agent policy configuration under governance controls.	Only authorized admins can update global model settings; changes are versioned and auditable.

11. Non-Functional Requirements

ID	Priority	Category	Requirement
NFR-SEC-001	P0	Security	All data in transit and at rest must be encrypted using enterprise-approved standards. Secrets and connector tokens must be stored in a managed secret store.
NFR-SEC-002	P0	Security	All retrieval, generation, and agent execution must enforce user permissions and source-level access constraints.
NFR-SEC-003	P0	Security	The system must support session timeout, remote logout, device revocation, and lost-device response for mobile.
NFR-AUD-001	P0	Auditability	Every generated artifact must retain provenance, generation time, model/config version where applicable, user context, and source references.
NFR-PRV-001	P0	Privacy	Recording, transcription, retention, redaction, and sharing behavior must be configurable by policy and visible to authorized users.
NFR-PERF-001	P1	Performance	Cached brief and decision list should load within 2 seconds under normal network conditions.
NFR-PERF-002	P1	Performance	Ask responses for normal questions should start streaming within 3 seconds when indexes are healthy; longer tasks must show progress.
NFR-PERF-003	P1	Performance	Live meeting transcript should appear with low enough latency to support in-meeting agent observations; exact target must be validated during pilot.
NFR-REL-001	P0	Reliability	Audio capture must continue through transient network loss and sync when the device reconnects, subject to mobile OS limits.
NFR-REL-002	P1	Reliability	Connector failures must not corrupt existing indexes; affected surfaces should show freshness and source health.
NFR-LOC-001	P1	Localization	English must be supported for v1. Urdu and bilingual transcript/brief support should be planned as a phased capability.
NFR-A11Y-001	P1	Accessibility	Mobile screens should meet enterprise accessibility expectations for readable typography, voice support, contrast, and touch targets.
NFR-OBS-001	P0	Observability	System health, connector health, agent failures, transcription failures, indexing lag, and generation errors must be observable by authorized operators.

12. Conceptual Data Model

The conceptual model below defines product entities. It is not a physical schema; engineering should convert it into a data architecture with source-system mapping, policy enforcement, versioning, and retention controls.

Entity	Description
Tenant	Workspace such as PM Office or a ministry cluster; owns policies, connectors, and roles.
User	Person using Mashwara with identity, role, ministry affiliation, delegation, and preferences.
Role / Policy	RBAC/ABAC rule defining what a user or agent can access and do.
Ministry / Entity	Government organization involved in decisions, meetings, and execution.
DataSource	Configured connector source such as Outlook mailbox, calendar, MoM repository, document store, or execution system.
SourceDocument	Normalized source object with provenance, permissions, content, metadata, and parse status.
EmailThread	Thread, messages, attachments, participants, reply state, and extracted signals.
CalendarEvent	Scheduled meeting metadata, participants, agenda, linked docs, and prep/summary status.
Meeting	Scheduled or ad-hoc meeting record with prep pack, live session, summary, decisions, and actions.
LiveSession	Mobile/room capture session with audio, transcript segments, sensitivity, and recording events.
TranscriptSegment	Timestamped transcript text, speaker label, confidence, audio pointer, and sensitivity flags.
Decision	Structured issue requiring approval, rejection, clarification, conditions, monitoring, or closure.
Recommendation	AI-generated recommended action with confidence, criteria, evidence, and alternatives.
Directive	Human-approved instruction or communication generated from a decision.
Receipt	Registered proof that a decision/directive was issued, with timestamp, reference, and recipients.
ActionItem	Execution task with owner, due date, status, source, evidence requirements, and escalation state.
ExecutionLog	Timeline of updates, evidence, comments, delays, escalations, and closure.
Agent	User-created or template-based routine with objective, scope, triggers, policy, version, and output routing.
AgentRun	Execution instance with context, source scope, observations, errors, and outputs.
AgentOutput	Signal, alert, report, suggested question, draft action, or decision candidate generated by an agent.
ConflictSignal	Detected disagreement across ministries, documents, meetings, or statements with evidence.
ConversationThread	Ask Mashwara history with messages, context, sources, and retention policy.
AuditEvent	Immutable event for access, generation, edit, send, export, policy change, or agent/recording activity.

12.1 Key Relationships

- User belongs to one or more roles and may be associated with one or more ministries/entities.
- DataSource produces SourceDocuments, EmailThreads, CalendarEvents, Meeting records, TranscriptSegments, and ExecutionLogs.
- Meeting can contain a LiveSession, PrepPack, PostSummary, Decisions, ActionItems, AgentRuns, and SourceDocuments.
- Decision links to Recommendation, EvidenceItems, Directive, Receipt, ActionItems, RelatedMeetings, RelatedDecisions, and AuditEvents.
- Agent contains configuration and versions; AgentRun produces AgentOutputs routed to Brief, Decision, Execution, Ask, or Notifications.
- Execution item links back to its originating decision, meeting, directive, or agent output and accumulates ExecutionLogs until closure.

13. Key User Journeys

Morning national brief

1. User opens Mashwara or receives secure morning notification.
2. Brief loads with critical decisions, execution delays, conflict signals, high-stakes meetings, and agent alerts.
3. User plays audio brief or taps into a decision/meeting/execution item.
4. Mashwara carries source context into the opened detail screen and Ask Mashwara suggestions.

Decision from email or MoM to directive

5. RAG ingestion identifies a decision candidate from email/MoM/source document.
6. System groups related sources, deduplicates similar items, and creates a decision in Needs Review or Ready lane.
7. Authorized official reviews recommendation, confidence, lineage, ministry positions, and source pack.
8. Official chooses action, edits directive, confirms send/registration, and receives receipt.

9. Execution item is created if the decision requires implementation.

Start ad-hoc live meeting on mobile

10. Authorized user taps Start Meeting.
11. User selects scheduled meeting or creates ad-hoc session, sets sensitivity, confirms consent policy, and starts recording.
12. Transcript appears live or near-live; detected decisions/actions and agent observations populate the live panel.
13. User ends meeting; Mashwara creates draft MoM, decisions, actions, and agent outputs for review.
14. Reviewer publishes accepted artifacts into Decisions, Meetings, and Execution.

Create and run an agent

15. User opens Agent Studio and describes objective in natural language.
16. Mashwara proposes agent configuration: scope, sources, triggers, output type, recipients, and sensitivity.
17. User reviews policy warnings, optionally test-runs the agent, then activates it.
18. Agent outputs flow to Brief, Agent Inbox, Decisions, Execution, or Notifications based on configuration.

Run agent during a live meeting

19. Meeting lead starts or opens a live session.
20. Lead selects one or more agents such as Commitment Tracker or Conflict Detector.
21. Agents monitor transcript segments and source context while recording continues.
22. Agent observations show timestamp, confidence, and evidence; accepted observations become actions, questions, or signals.

Execution delay escalation

23. Execution Monitor detects overdue or no-update action.
24. Item appears in Brief and Execution dashboard with owner, ministry, source, deadline, and delay reason if known.
25. Authorized user requests update, escalates, or asks Mashwara to summarize blocker history.
26. Updated status and evidence are logged; closure requires evidence or authorized override.

Ask institutional memory question

27. User asks a question such as: What did we decide about this policy in the last three meetings?
28. Mashwara retrieves authorized emails, MoMs, transcripts, decisions, and execution logs.
29. Answer includes concise synthesis, confidence, uncertainty, and source links.
30. User can turn answer into an agent, decision candidate, memo, or follow-up question when authorized.

14. Enterprise Rollout Plan

Phase	Scope	Notes
Phase 0 - Enterprise readiness	Identity, RBAC/ABAC, connector approvals, retention policy, audit logs, deployment model, and data classification policy.	Required before any real official data is ingested.
Phase 1 - RAG foundation pilot	Outlook RAG, Calendar RAG, MoM RAG, National Brief, Decision Engine, meeting prep/post summaries, Ask Mashwara.	Run with one ministry or controlled PM Office unit.
Phase 2 - Execution monitor	Decision-to-action lifecycle, owner assignment, status, overdue detection, silent failure detection, executive dashboards.	Use with secretariat/focal point operating model.
Phase 3 - Live meeting capture	Mobile Start Meeting, recording policy, live transcription, draft MoM, detected decisions/actions, transcript evidence.	Start with internal review or secretary-level meetings before highly sensitive meetings.
Phase 4 - User-created agents	Natural language Agent Studio, templates, live meeting agents, historical RAG agents, agent governance and outputs.	Enable gradually with policy guardrails and template review.
Phase 5 - Cross-ministry scale	Cross-ministry conflict detection, PM-level execution rollups, shared agents, enterprise analytics, expanded connectors.	Requires mature data governance and operating procedures.

15. Success Metrics

Metric	Target / Measurement
Brief adoption	80% of target users open or listen to brief on working days during pilot.
Brief speed	Executives can complete morning brief loop in under 5 minutes for normal days.

Metric	Target / Measurement
Decision explainability	100% of recommendations show source evidence, confidence, and reasoning.
Decision extraction quality	Decision candidate extraction precision/recall measured against human-labeled sample; thresholds set during pilot.
Live meeting review time	Draft MoM and extracted actions are available for review within minutes of meeting end, subject to transcript processing.
Agent creation usability	Authorized users can create a valid agent from natural language in under 3 minutes for template-compatible use cases.
Execution visibility	Overdue and silent-failure items are visible in Brief/Execution without manual spreadsheet tracking.
Permission safety	Zero known incidents where a user or agent receives unauthorized content.
Source coverage	Connector health dashboard reports ingestion freshness and failures for all approved sources.

16. Risks, Assumptions, and Open Questions

16.1 Key Risks and Mitigations

Risk	Description	Mitigation
Recording sensitivity	Officials may resist recording or some meetings may be legally/politically sensitive.	Make recording policy configurable; support off-record segments; begin with internal/secretariat meetings; require explicit human review.
Data quality	Emails, MoMs, and documents may be inconsistent, duplicated, scanned, or incomplete.	Add normalization, human review, source health dashboards, correction feedback, and confidence indicators.
Permission leakage	Cross-source RAG and agents can accidentally expose restricted data if not policy-scoped.	Enforce source-level permissions before retrieval/generation; audit all runs; run adversarial access tests.
Over-trust in AI	Users may treat AI recommendations as final.	Require human approval for directives; show confidence and evidence; separate drafts from official records.
Agent sprawl	User-created agents may become noisy or redundant.	Use templates, ownership, schedules, quotas, deactivation, run history, and admin governance.
Operational burden	Connector failures and review queues may require a new operating model.	Define private office/secretariat reviewer roles and admin SLAs before pilot.
Language and transcription	Mixed English/Urdu/regional language conversations may reduce transcript quality.	Pilot language support, speaker correction, and confidence displays; phase bilingual capabilities.

16.2 Working Assumptions

- The first pilot will use approved official data sources and a controlled set of users.
- RAG outputs, meeting summaries, agent outputs, and decisions are considered drafts until reviewed where policy requires.
- Sensitive recording policies must be agreed before enabling live meeting capture for senior meetings.
- Directives are always human-approved before dispatch or registration.
- The product must run with English support in v1; Urdu/bilingual support is planned but depends on policy and transcription validation.
- Enterprise admin and security functions are required even if the first user-facing pilot is mobile-first.

16.3 Open Questions

- What is the authoritative identity provider and device management model for PM/Minister mobile access?
- Which mailboxes, calendars, MoM repositories, document stores, and execution systems are approved for Phase 1?
- What is the policy for recording consent, off-record segments, sensitive meetings, and transcript retention?
- Who can create agents, who can approve shared agents, and what output types require human review?
- What is the official format for Pakistan decision/directive references and receipts?
- Should v1 support English only, Urdu only, or bilingual summaries/transcripts?
- Which roles can see cross-ministry information, and what delegation rules apply when staff support an executive?
- What is the legal records status of AI-generated MoMs, summaries, and transcripts before human approval?
- What are the target latency and accuracy thresholds for live transcription under real meeting conditions?
- What deployment model is required: sovereign cloud, government cloud, private cloud, or on-premise?

17. Out of Scope

- Autonomous final decisions or autonomous dispatch of directives without human confirmation.
- Public citizen engagement, media communications, or press release automation.
- Calendar editing, scheduling negotiation, or meeting cancellation workflows.
- Provincial government integration in v1 unless explicitly included in the pilot.
- Classified or national-security-sensitive meetings until recording, retention, and clearance controls are formally approved.
- Full case-management/workflow replacement for existing government systems.
- Political campaign, party, or non-official communication use cases.
- External open-web news/social intelligence unless separately approved as a future source layer.

Appendix A. Requirement Index

ID	Priority	Group	Requirement
FR-APP-001	P0	App Shell, Identity, and Navigation	Authenticate users through enterprise SSO with MFA and optional biometric session unlock on mobile.
FR-APP-002	P0	App Shell, Identity, and Navigation	Render navigation based on role and authorized modules: Brief, Decisions, Meetings, Live, Agents, Execution, Ask, and Settings/Admin as applicable.
FR-APP-003	P0	App Shell, Identity, and Navigation	Expose Ask Mashwara globally as a persistent, context-aware control across core user screens.
FR-APP-004	P1	App Shell, Identity, and Navigation	Support secure push notifications for prep packs, meeting summaries, high-priority decisions, agent alerts, and overdue execution items.
FR-APP-005	P1	App Shell, Identity, and Navigation	Allow user preferences for brief time, notification intensity, default language, audio voice, and sensitive-mode defaults.
FR-APP-006	P1	App Shell, Identity, and Navigation	Provide offline read access to the latest downloaded brief, meeting prep pack, and cached decision detail, excluding send/execute actions.
FR-RAG-001	P0	Data Ingestion and RAG Foundation	Ingest Outlook email threads and attachments from approved mailboxes and distribution groups.
FR-RAG-002	P0	Data Ingestion and RAG Foundation	Ingest official calendars for authorized PM/Minister/Secretary accounts and selected shared calendars.
FR-RAG-003	P0	Data Ingestion and RAG Foundation	Ingest historical minutes of meeting repositories and formal decision records.
FR-RAG-004	P0	Data Ingestion and RAG Foundation	Ingest upcoming meeting agendas and pre-read material before meeting time.
FR-RAG-005	P0	Data Ingestion and RAG Foundation	Normalize and classify unstructured content into source document, email thread, meeting, decision candidate, action item, stakeholder position, and evidence objects.
FR-RAG-006	P0	Data Ingestion and RAG Foundation	Apply permission-aware retrieval for every RAG response, brief item, decision, agent output, and search result.
FR-RAG-007	P0	Data Ingestion and RAG Foundation	Maintain durable provenance for every generated insight.
FR-RAG-008	P1	Data Ingestion and RAG Foundation	Detect duplicates, near-duplicates, thread continuations, and renamed documents across sources.
FR-RAG-009	P1	Data Ingestion and RAG Foundation	Capture human corrections to improve future extraction and classification.
FR-BRF-001	P0	National Brief	Generate a daily National Brief for each executive based on authorized sources and role scope.
FR-BRF-002	P0	National Brief	Include required sections: Critical Decisions Today, Pending Since Last Cabinet/Committee, Execution Delays, Cross-Ministry Conflicts, Upcoming High-Stakes Meetings, and Agent Alerts.
FR-BRF-003	P0	National Brief	Support a 3-5 minute audio brief with chapter navigation.
FR-BRF-004	P0	National Brief	Make every brief item drillable into the underlying decision, meeting, source pack, execution item, or agent output.
FR-BRF-005	P1	National Brief	Highlight carryovers and stale items based on prior meetings, unresolved decisions, and execution status.
FR-BRF-006	P1	National Brief	Allow the private office or authorized reviewer to mark brief items as sensitive, suppress from audio, or require review before publication.
FR-DEC-001	P0	Decision Engine	Automatically extract decision candidates from emails, meetings, documents, live transcripts, and agent outputs.
FR-DEC-002	P0	Decision Engine	Support decision lanes: Needs Review, Ready, Waiting on Entity, Executing, Deferred, Completed, and Closed Without Action.
FR-DEC-003	P0	Decision Engine	Display decision detail with at-a-glance summary, recommendation, confidence, ministries involved, stakeholder positions, source pack, related meetings, related decisions, and execution status.
FR-DEC-004	P0	Decision Engine	Generate one of four primary recommendations: Approve, Approve with Conditions, Return for Clarification, or Reject.
FR-DEC-005	P0	Decision Engine	Provide a one-tap action path into an editable directive draft.
FR-DEC-006	P0	Decision Engine	Require explicit human confirmation before sending or registering any directive.
FR-DEC-007	P0	Decision Engine	Generate a registered decision receipt after send/approval with reference, timestamp, action, recipients, and status row.
FR-DEC-008	P0	Decision Engine	Maintain decision lineage from first mention to current state.
FR-DEC-009	P1	Decision Engine	Detect potential external completion from email replies, signed documents, or execution evidence.

ID	Priority	Group	Requirement
FR-DEC-010	P0	Decision Engine	Maintain a full audit trail of decision state, source changes, reviewer edits, approvals, deferrals, and dispatch events.
FR-DEC-011	P1	Decision Engine	Detect conflicting ministry positions and attach a conflict signal to the decision.
FR-DEC-012	P2	Decision Engine	Suggest precedent decisions and similar historical cases.
FR-MTG-001	P0	Meetings Intelligence	Show a role-aware meeting timeline with scheduled, in-progress, completed, and ad-hoc live sessions.
FR-MTG-002	P0	Meetings Intelligence	Generate a prep pack for each upcoming meeting with context, agenda, issue history, stakeholders, prior decisions, unresolved actions, documents, and suggested talking points.
FR-MTG-003	P0	Meetings Intelligence	Generate post-meeting summaries containing summary, decisions, action items, commitments, follow-ups, owners, due dates, and source evidence.
FR-MTG-004	P0	Meetings Intelligence	Provide a continuity layer for recurring or related meetings.
FR-MTG-005	P1	Meetings Intelligence	Support in-app document preview for approved formats and source packs.
FR-MTG-006	P1	Meetings Intelligence	Send configurable prep and summary notifications.
FR-MTG-007	P0	Meetings Intelligence	Link meeting outcomes to Decision Engine and Execution Monitor.
FR-LIVE-001	P0	Live Meeting Capture	Allow authorized users to tap Start Meeting on mobile for scheduled or ad-hoc meetings.
FR-LIVE-002	P0	Live Meeting Capture	Allow the user to attach the live session to a calendar meeting or create an ad-hoc meeting record.
FR-LIVE-003	P0	Live Meeting Capture	Enforce configurable recording consent and sensitivity policy before capture begins.
FR-LIVE-004	P0	Live Meeting Capture	Capture mobile audio reliably with clear recording status and network/offline state.
FR-LIVE-005	P0	Live Meeting Capture	Generate live or near-live transcription with speaker segmentation where supported.
FR-LIVE-006	P0	Live Meeting Capture	Detect decisions, commitments, action items, risks, disagreements, deadlines, and named owners during the meeting.
FR-LIVE-007	P0	Live Meeting Capture	Allow users to run selected agents on top of the live meeting.
FR-LIVE-008	P1	Live Meeting Capture	Support in-meeting controls: pause recording, mark off-record segment, bookmark moment, tag topic, and flag sensitive segment.
FR-LIVE-009	P0	Live Meeting Capture	Generate a review package when the meeting ends.
FR-LIVE-010	P0	Live Meeting Capture	Route accepted decisions and actions into Decision Engine and Execution Monitor.
FR-LIVE-011	P1	Live Meeting Capture	Allow post-meeting correction of transcript, speaker labels, summaries, owners, and deadlines.
FR-LIVE-012	P2	Live Meeting Capture	Support multi-device or room-system capture when enterprise hardware is available.
FR-AGT-001	P0	Agent Studio and Agent Execution	Allow authorized users to create agents using natural language.
FR-AGT-002	P0	Agent Studio and Agent Execution	Provide approved agent templates for monitoring, meeting capture, decision review, conflict detection, execution tracking, and weekly reporting.
FR-AGT-003	P0	Agent Studio and Agent Execution	Allow agent configuration of data scope, triggers, schedule, output type, recipients, and sensitivity.
FR-AGT-004	P0	Agent Studio and Agent Execution	Allow agents to run on historical RAG sources, upcoming meetings, a selected decision, a live meeting, or execution logs.
FR-AGT-005	P0	Agent Studio and Agent Execution	Support real-time meeting agents that detect commitments, contradictions, missing owners, policy references, and follow-up questions.
FR-AGT-006	P0	Agent Studio and Agent Execution	Support agent output types: signal, alert, report, suggested question, draft action, and decision candidate.
FR-AGT-007	P0	Agent Studio and Agent Execution	Allow users to pause, resume, run now, edit, duplicate, deactivate, and view run history for their agents.
FR-AGT-008	P0	Agent Studio and Agent Execution	Apply guardrails: agents cannot access unauthorized sources, cannot dispatch directives autonomously, and cannot override sensitivity policy.
FR-AGT-009	P1	Agent Studio and Agent Execution	Provide test-run preview before activating an agent.
FR-AGT-010	P1	Agent Studio and Agent Execution	Support shared agents managed by a private office or ministry team.
FR-AGT-011	P2	Agent Studio and Agent Execution	Provide an enterprise template library governed by admins.
FR-EXE-001	P0	Execution Monitor	Create execution items from directives, MoMs, live meeting actions, and accepted agent outputs.
FR-EXE-002	P0	Execution Monitor	Track status across Not Started, In Progress, At Risk, Blocked, Completed, Deferred, and Cancelled.
FR-EXE-003	P0	Execution Monitor	Detect overdue actions and silent failures.
FR-EXE-004	P1	Execution Monitor	Support escalation from ministry focal point to Secretary, Minister, or PM Office based on policy.
FR-EXE-005	P0	Execution Monitor	Allow closure only with completion evidence or authorized override.
FR-EXE-006	P1	Execution Monitor	Provide ministry and cross-ministry execution dashboards.
FR-EXE-007	P1	Execution Monitor	Generate weekly execution summaries for authorized leaders.
FR-EXE-008	P2	Execution Monitor	Support integration with external workflow or document-signature systems where approved.
FR-ASK-001	P0	Ask Mashwara	Allow users to query across authorized emails, meetings, MoMs, transcripts, decisions, agents, documents, and execution logs.
FR-ASK-002	P0	Ask Mashwara	Show source citations and confidence for every factual answer.
FR-ASK-003	P0	Ask Mashwara	Clearly state uncertainty, missing data, or insufficient authorization.
FR-ASK-004	P0	Ask Mashwara	Provide context-aware suggestion chips on key screens.
FR-ASK-005	P1	Ask Mashwara	Support voice input and voice playback where policy allows.
FR-ASK-006	P1	Ask Mashwara	Allow users to upload or reference a file for evaluation.
FR-ASK-007	P1	Ask Mashwara	Support conversation history with topic, source scope, and continuation.
FR-ASK-008	P0	Ask Mashwara	Block or safely decline requests that would violate role permissions, data policy, or autonomous decision boundaries.
FR-ADM-001	P0	Enterprise Administration and Governance	Provide tenant and workspace configuration for PM Office and ministries.

ID	Priority	Group	Requirement
FR-ADM-002	P0	Enterprise Administration and Governance	Provide RBAC/ABAC policy management.
FR-ADM-003	P0	Enterprise Administration and Governance	Provide connector management for Outlook, calendars, MoM repositories, document repositories, and execution systems.
FR-ADM-004	P0	Enterprise Administration and Governance	Provide retention and deletion policy controls for transcripts, audio, summaries, agent outputs, decisions, and audit logs.
FR-ADM-005	P0	Enterprise Administration and Governance	Provide audit logs for access, generation, exports, edits, approvals, agent runs, recording events, and policy changes.
FR-ADM-006	P0	Enterprise Administration and Governance	Provide recording consent and sensitivity policy configuration.
FR-ADM-007	P1	Enterprise Administration and Governance	Provide data quality and ingestion health dashboards.
FR-ADM-008	P1	Enterprise Administration and Governance	Provide model, prompt, and agent policy configuration under governance controls.

End of PRD